

Contents

Solution Guide: Intelligent App Development with Copilot Stack Scenario	2
Step 1: Deploy Azure OpenAI Service and LLM Models	2
Task 1: Deploy an Azure OpenAI Service	2
Task 2: Deploy LLM models in Azure OpenAI Studio	3
Step 2: Explore Semantic Kernel – Solution Guide	7
Task 1: Clone the repository for this course	7
Task 2: Retrieving the Azure OpenAI Service values	7
Task 3: Run Jupyter Notebooks to get started with Semantic Kernel in the Python programming language	9
Optional Tasks	14
Step 3: Run the Chat Copilot App Locally - Solution Guide	16
Task 1: Retrieving the Azure OpenAI Service values	16
Task 2: Cloning the Chat-Copilot GitHub Repo	18
Task 3: Setting up the Environment	19
Task 4: Configure and run the Chat Copilot App locally	20

Solution Guide: Intelligent App Development with Copilot Stack Scenario

Step 1: Deploy Azure OpenAI Service and LLM Models

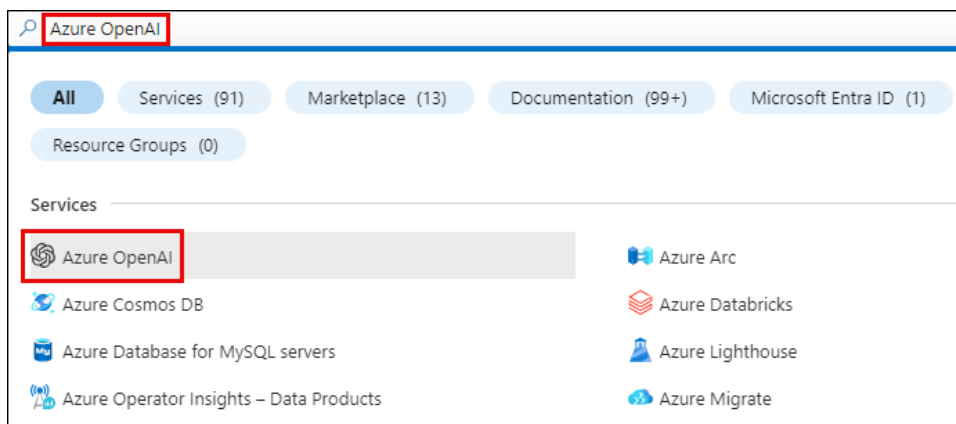
Notes & Guidance

Task 01 is all about helping the student set up the prerequisites for this Step. This includes necessary installations, environment options, and other libraries needed.

Task 1: Deploy an Azure OpenAI Service

In this task, you'll be deploying an Azure OpenAI service in the Azure Portal.

1. In the Azure Portal, search for **Azure OpenAI** and select it.



2. On **Azure AI Services | Azure OpenAI** blade, click on **create** enter the details required and deploy the Azure Open AI service.

Create Azure OpenAI

1 Basics 2 Network 3 Tags 4 Review + submit

Enable new business solutions with OpenAI's language generation capabilities powered by GPT-3 models. These models have been pretrained with trillions of words and can easily adapt to your scenario with a few short examples provided at inference. Apply them to numerous scenarios, from summarization to content and code generation.

[Learn more](#)

Project Details

Subscription * ⓘ Azure Lab: ▼

Resource group * ⓘ Activate- ▼
[Create new](#)

Instance Details

Region ⓘ East US ▼

Name * ⓘ OpenAI-1180344 ✓

Pricing tier * ⓘ Standard S0 ▼

[View full pricing details](#)

Content review policy

To detect and mitigate harmful use of the Azure OpenAI Service, Microsoft logs the content you send to the Completions and image generations APIs as well as the content it sends back. If content is flagged by the service's filters, it may be reviewed by a Microsoft full-time employee.

[Learn more about how Microsoft processes, uses, and stores your data](#)

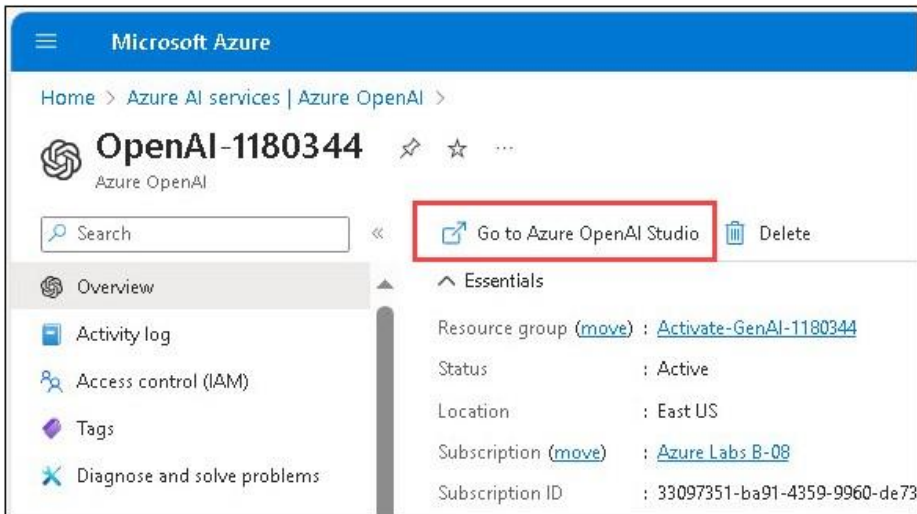
[Apply for modified content filters and abuse monitoring](#)

Previous Next

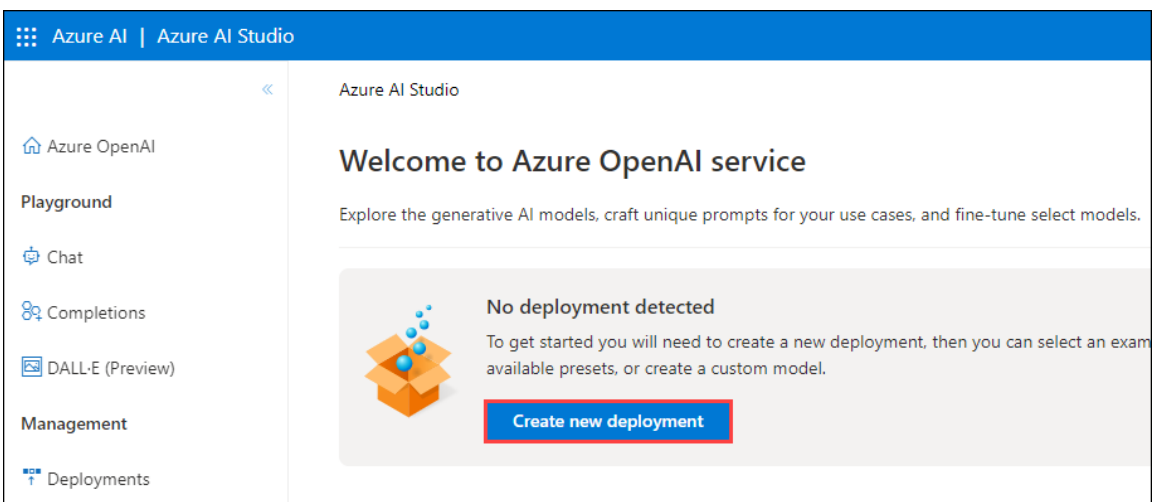
Task 2: Deploy LLM models in Azure OpenAI Studio

In this task, you'll be deploying the OpenAI models from Azure OpenAI Studio.

1. In the Azure OpenAI resource pane, click on **Go to Azure OpenAI Studio**; it will navigate to **Azure AI Studio**.



2. On the **Welcome to Azure OpenAI Service** page, click on **Create new deployment**.



3. In the **Deployments** page, click on **+ Create new deployment**.
4. Within the **Deploy model** pop-up interface, enter the following details and then click on **Advanced options (3)**, followed by scaling down the **Tokens per Minute Rate Limit (thousands) (4)**:
 - **Select a model:** gpt-35-turbo (1)
 - **Model version:** *Use the default version* (2)
 - **Deployment name:** gpt-35-turbo
 - **Tokens per Minute Rate Limit (thousands):** 20K

Set up a deployment to make API calls against a provided base model or a custom model. Finished deployments are available for use. Your deployment status will move to succeeded when the deployment is complete and ready for use.

Select a model ⓘ **1**

gpt-35-turbo

Model version ⓘ **2**

Auto-update to default

Deployment name ⓘ

text-turbo

3 ⚙️ Advanced options ▾

Content Filter ⓘ

Default

10000K tokens per minute quota available for your deployment

We were unable to find the exact quota available as the call to the usages API failed. This could be due to insufficient permissions.

Tokens per Minute Rate Limit (thousands) ⓘ **4**

0 1K

Corresponding requests per minute (RPM) = 6

Enable Dynamic Quota ⓘ

☒ Enabled

Create Cancel

5. Click on the **Create** button to deploy a model that you will be playing around with as you proceed.
6. In the **Deployments** page again, click on + **Create new deployment**.
7. Within the **Deploy model** pop-up interface, enter the following details and then click on **Advanced options (3)**, followed by scaling down the **Tokens per Minute Rate Limit (thousands) (4)**:
 - **Select a model:** text-embedding-ada-002 (1)
 - **Model version:** *Use the default version* (2)
 - **Deployment name:** text-embedding-ada-002
 - **Tokens per Minute Rate Limit (thousands):** 20K

Deploy model

Set up a deployment to make API calls against a provided base model or a custom model. Finished deployments are available for use. Your deployment status will move to succeeded when the deployment is complete and ready for use.

Select a model ⓘ

text-embedding-ada-002

Model version ⓘ

2 (Default)

Deployment name ⓘ

text-ada-002

⚙️ Advanced options ▾

Content Filter ⓘ

Default

ⓘ 240K tokens per minute quota available for your deployment

Tokens per Minute Rate Limit (thousands) ⓘ

0 1K

Corresponding requests per minute (RPM) = 6

Enable Dynamic Quota ⓘ

☒ Enabled

Create Cancel

- Click on the **Create** button to deploy a model that you will be playing around with as you proceed.

Step 2: Explore Semantic Kernel – Solution Guide

Task 1: Clone the repository for this course

If you have not already cloned the **Semantic Kernel GitHub Repo** code repository to the environment where you're working on this lab, follow these steps to do so. Otherwise, open the cloned folder in Visual Studio Code.

1. Open Visual Studio Code.
2. Create and Navigate to the directory `C:/Users/azureuser`.

```
...
```

```
cd C:/Users/azureuser
```

```
...
```

3. Clone the given repository.

```
...
```

```
git clone https://github.com/microsoft/semantic-kernel
```

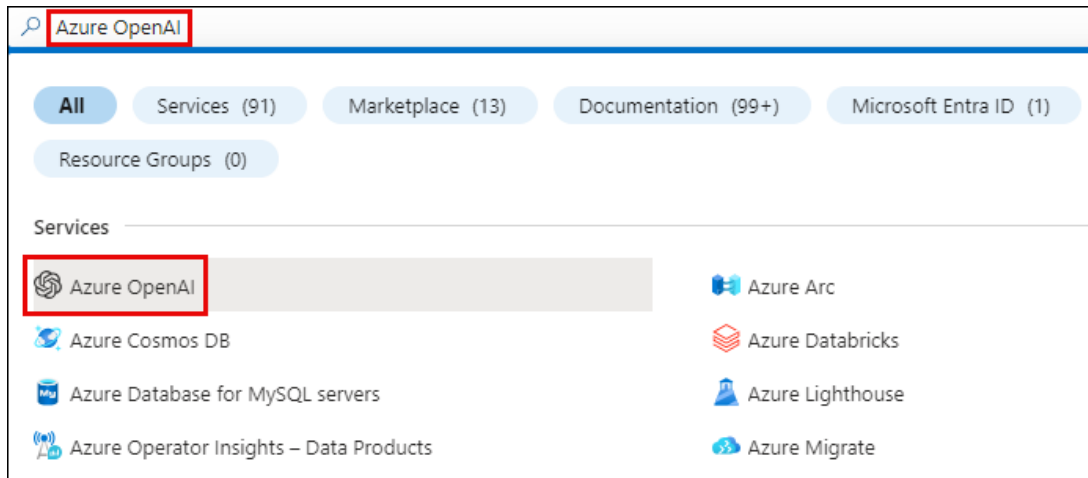
```
...
```

4. When the repository has been cloned, open the folder in Visual Studio Code.

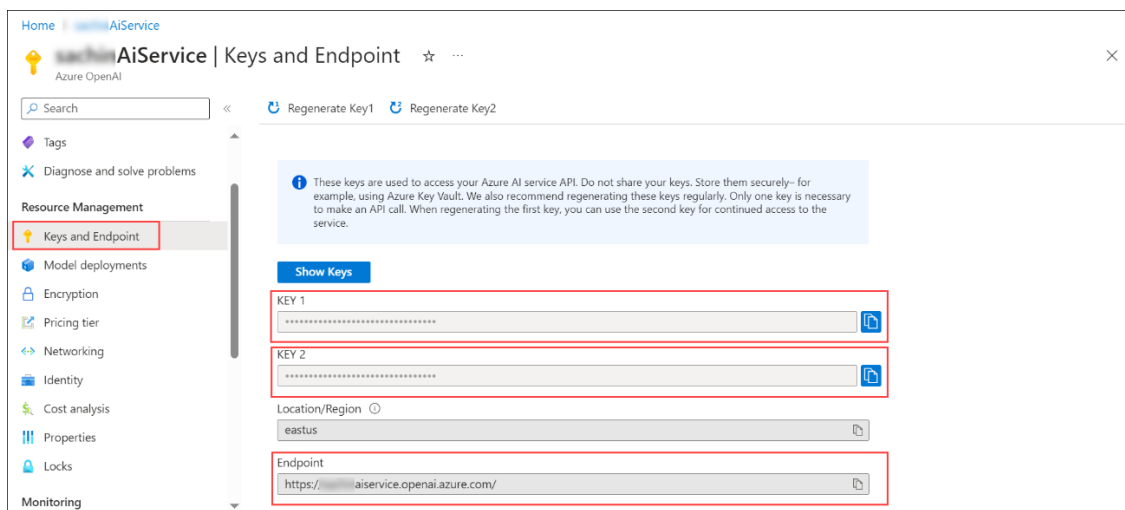
Task 2: Retrieving the Azure OpenAI Service values

From the Azure portal, you need to retrieve the Azure OpenAI Service Key, Endpoint, and LLM model name deployed in the previous Step.

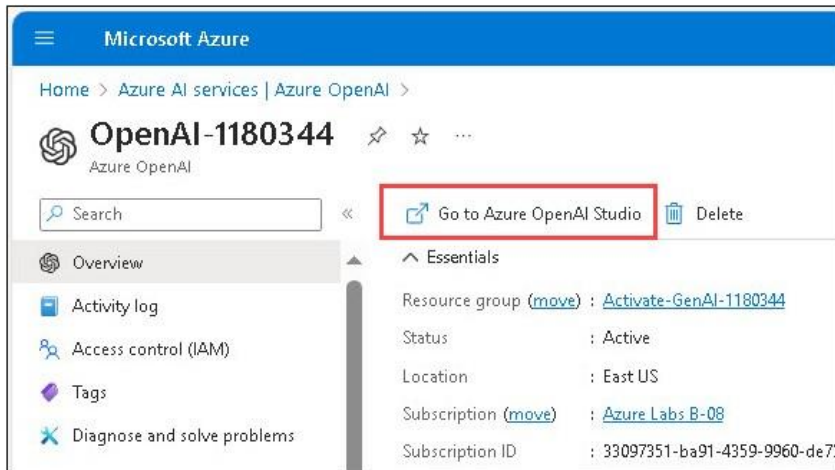
1. In the Azure Portal, search for **Azure OpenAI** and select it.



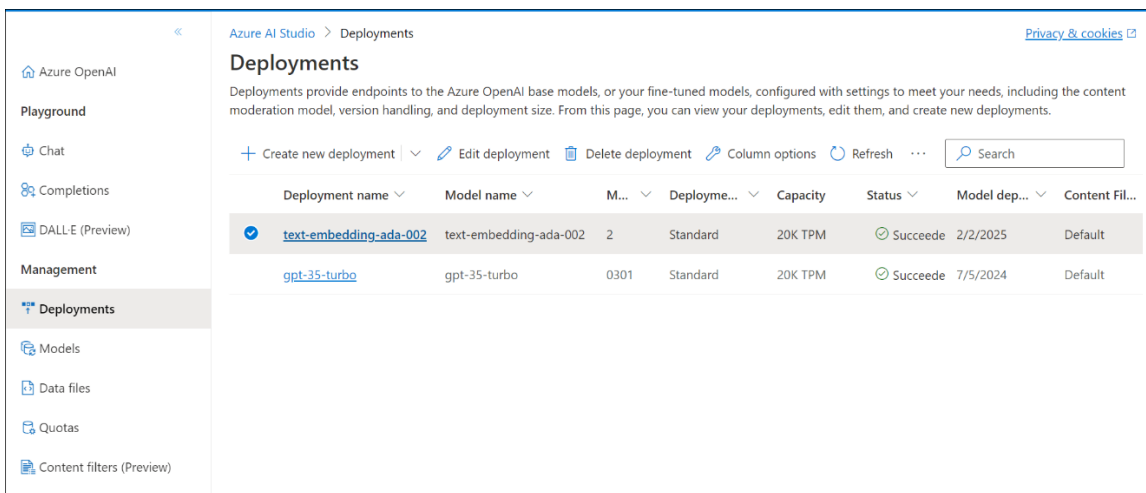
2. Go to the Azure OpenAI resource that you created previously and choose **Keys and Endpoints** from the left pane.
3. Copy any of the 2 Keys and the Endpoint and store it somewhere for future reference.



4. From the Overview page, click **Go to Azure OpenAI Studio** to go to your deployed models.



5. Go to **Deployments**, and you will see your deployed AI models. Copy the deployment names of your AI model that you deployed previously and store them somewhere for future reference.



Task 3: Run Jupyter Notebooks to get started with Semantic Kernel in the Python programming language

Inside VS Code, you need to run the required notebooks cell-by-cell successfully and observe the outputs for better understanding.

1. Go to the cloned folder in your **VS Code**, navigate to **semantic-kernel/python/samples/getting_starded**, and create a new **.env** file.
2. Go to **.env.example** , copy all its contents, and paste them into the newly created **.env** file.

The screenshot shows the VS Code Explorer on the left with the file tree expanded to the `notebooks` directory. The `.env.example` file is selected. The main editor area displays the content of `.env.example`, which is a template for environment variables for the Semantic Kernel Python notebooks.

```
semantic-kernel > python > notebooks > $ .env.example
1 OPENAI_API_KEY=""
2 OPENAI_ORG_ID=""
3 AZURE_OPENAI_DEPLOYMENT_NAME=""
4 AZURE_OPENAI_ENDPOINT=""
5 AZURE_OPENAI_API_KEY=""
6 AZURE_AISEARCH_API_KEY=""
7 AZURE_AISEARCH_URL=""
8
```

3. Modify the `.env` file by updating the **Azure OpenAI Model name, Endpoint, and Key** that you have retrieved in the earlier steps.

The screenshot shows the VS Code Explorer on the left with the file tree expanded to the `notebooks` directory. The `.env` file is selected. The main editor area displays the content of `.env`, which has been updated with values from the previous steps.

```
semantic-kernel > python > notebooks > $ .env
1 OPENAI_API_KEY=""
2 OPENAI_ORG_ID=""
3 AZURE_OPENAI_DEPLOYMENT_NAME="gpt-35-turbo"
4 AZURE_OPENAI_ENDPOINT="https://[redacted].openai.azure.com/"
5 AZURE_OPENAI_API_KEY="[redacted]"
6 AZURE_AISEARCH_API_KEY=""
7 AZURE_AISEARCH_URL=""
8
```

4. Navigate to `00-getting-started.ipynb` and run the notebook cell-by-cell, where you will be importing Semantic Kernel SDK.
Note: Skip the cell of codes pertaining to OpenAI, i.e., **Option 1** in the notebook, and only run the cell of codes related to Azure **OpenAI**, i.e., starting from **Option 2** in the notebook.
5. Update the values of **deployment_name**, **endpoint**, and **api_key** in **Option 2**, which you have already retrieved, and run the function.

Option 2: using Azure OpenAI

Step 2: Add your Azure Open AI Service key settings to a `.env` file in the same folder:

```
AZURE_OPENAI_API_KEY="..."
AZURE_OPENAI_ENDPOINT="https://..."
AZURE_OPENAI_DEPLOYMENT_NAME="..."
```

Use "keyword arguments" to instantiate an Azure OpenAI Chat Completion service and add it to the kernel:

```
from semantic_kernel.connectors.ai.open_ai import AzureChatCompletion

deployment, api_key, endpoint = sk.azure_openai_settings_from_dot_env()


kernel.add_chat_service("chat_completion", AzureChatCompletion(deployment_name="gpt-35-turbo", endpoint="https://...aiservice.openai.azure.com/"))
```

[3] ✓ 1.4s Python

... <semantic_kernel.kernel.Kernel at 0x29ef5740150>

Run the remaining code as usual.

- After successfully completing the previous notebook, navigate to `01-basic-loading-the-kernel.ipynb` and run the notebook cell-by-cell. Choose **Python <your_version>** whenever a prompt appears on top of the screen.



- Update the values of **deployment_name**, **endpoint**, and **api_key**, which you noted earlier, in the following cell of code before executing it.

EXPLORER

00-getting-started.ipynb • 01-basic-loading-the-kernel.ipynb • 03-semantic-function-inline.ipynb

01-basic-loading-the-kernel.ipynb

Basic Loading of the Kernel

Python 3.11.4

The SDK currently supports OpenAI and Azure OpenAI, other services will be added over time.

If you need an Azure OpenAI key, go [here](#).

```
kernel = sk.Kernel()

kernel.add_chat_service(
    "Azure_curie",  # We are adding a text service
    AzureChatCompletion(  # The alias we can use in prompt templates' config.json
        deployment_name="gpt-35-turbo",  # Azure OpenAI "Deployment name"
        endpoint="https://...aiservice.openai.azure.com/",  # Azure OpenAI "Endpoint"
        api_key="...4821be9ff8ccb9b62"  # Azure OpenAI "Key"
    )
)

kernel.add_chat_service(
    "OpenAI_chat_gpt",  # We are adding a text service
    OpenAIChatCompletion(  # The alias we can use in prompt templates' config.json
        ai_model_id="gpt-3.5-turbo",  # OpenAI Model Name
        api_key="...your OpenAI API Key...",  # OpenAI API key
        org_id="...your OpenAI Org ID..."  # "optional" OpenAI Organization ID
    )
)
```

[3] ✓ 0.1s Python

... <semantic_kernel.kernel.Kernel at 0x2be14af84d0>

- After successfully running the previous notebook, navigate to `02-running-prompts-from-file.ipynb` and run the notebook cell-by-cell. Choose **Python <your_version>** whenever a prompt appears on top of the screen.
- In the following function, modify **useAzureOpenAI** to **True**, set the values of **deployment_name**, **endpoint**, and **api_key**, which you noted earlier, and execute it.

Run the remaining cells of code as usual.

- After successfully completing the previous notebook, navigate to `03-semantic-function-inline.ipynb` and run the notebook cell-by-cell. Choose **Python <your_version>** whenever a prompt appears on top of the screen. Here, you need an OpenAI model that supports text completion. For that, you need to navigate to Azure OpenAI Studio and deploy a `text-davinci-003` model with a TPM capacity of 10K.
- In the following function, set the values of **deployment_name**, **endpoint**, and **api_key** for your **text-davinci-003** model, which you noted earlier, and execute it.

```
import semantic_kernel as sk
from semantic_kernel.connectors.ai.open_ai import AzureTextCompletion, OpenAITextCompletion

kernel = sk.Kernel()

useAzureOpenAI = False

# Configure AI service used by the kernel
if useAzureOpenAI:
    deployment, api_key, endpoint = sk.azure_openai_settings_from_dot_env()
    azure_text_service = AzureTextCompletion(deployment_name="text", endpoint=endpoint, api_key=api_key) # set the deployment name to the value of deployment_name
    kernel.add_text_completion_service("dv", azure_text_service)
else:
    api_key, org_id = sk.openai_settings_from_dot_env()
    oai_text_service = OpenAITextCompletion(ai_model_id="text-davinci-003", api_key=api_key, org_id=org_id)
    kernel.add_text_completion_service("dv", oai_text_service)
```

Python

Run the remaining cells of code as usual.

12. In the following function, modify **useAzureOpenAI** to **True**, set the values of **deployment_name**, **endpoint**, and **api_key**, which you noted earlier, and execute it.

```
import semantic_kernel as sk
from semantic_kernel.connectors.ai.open_ai import AzureChatCompletion, OpenAIChatCompletion

kernel = sk.Kernel()

useAzureOpenAI = True

# Configure AI service used by the kernel
if useAzureOpenAI:
    deployment, api_key, endpoint = sk.azure_openai_settings_from_dot_env()
    azure_chat_service = AzureChatCompletion(deployment_name="gpt-35-turbo", endpoint="https://aiservice.openai.azure.com/", api_key=api_key)
    kernel.add_chat_service("chat_completion", azure_chat_service)
else:
    api_key, org_id = sk.openai_settings_from_dot_env()
    oai_chat_service = OpenAIChatCompletion(ai_model_id="gpt-3.5-turbo", api_key=api_key, org_id=org_id)
    kernel.add_chat_service("chat-gpt", oai_chat_service)
```

✓ 4.1s Python

Run the remaining cells of code as usual.

13. After successfully completing the previous notebook, navigate to 04-context-variables-chat.ipynb and run the notebook cell-by-cell. Choose **Python <your_version>** whenever a prompt appears on top of the screen.
14. In the following function, modify **useAzureOpenAI** to **True**, set the values of **deployment_name**, **endpoint**, and **api_key**, which you noted earlier, and execute it.

```
[1] ^ 132
    kernel.sendMessage("Bbf-3-2", obseuicmefcombjefjou(97-wogej-7q="Bbf-3-2-furpo", ebf-keλ=ebf-keλ, ouE-7q=ouE-7q))
    ebf-keλ, ouE-7q = ek.obseuicmefcombjefjou(97-wogej-7q="Bbf-3-2-furpo", ebf-keλ=ebf-keλ, ouE-7q=ouE-7q)
    736:
    kernel.sendMessage("cmef-combjefjou", vzinucmefcombjefjou(qebjolwewf-uwwea_Bbf-32-furpo, ewqbozurf="ufbz:\\ 97seuicmefcombjefjou.obseuicmefcombjefjou.co
    qebjolwewf, ebf-keλ, ewqbozurf = ek.vzinucmefcombjefjou(qebjolwewf-uwwea_Bbf-32-furpo, ewqbozurf="ufbz:\\ 97seuicmefcombjefjou.obseuicmefcombjefjou.co
    74 nzevzinucmefcombjefjou:
    75 coufzBne vj pscfkwq nzeq pl the kernel
    76vzinucmefcombjefjou = 136
    77
    78kernel = ek.kernel()
    79
    80from semantickernel.connectors.ai.obseuicmefcombjefjou obseuicmefcombjefjou, vzinucmefcombjefjou
    81import semantickernel as ek
```

Run the remaining cells of code as usual. An example of output appearing is like this:

```
Chatbot: You may learn about the world events' people' and there that shared ancient Greece' including the rise of democracy' the Persian Wars' the Pe
...
[10] ^ 032
    await chat("It I read that book' what exactly may I learn about Greek history?")

Chatbot: "The Greeks: A portrait of life and culture" is a comprehensive overview of ancient Greek culture' history' and mythology. It explores the Gre
...
[9] ^ 532
    await chat("What sounds interesting' what is it about?")

Chatbot: How about "The Greeks: A portrait of life and culture" by Paul Cartledge on "The Birth of Classical Europe: A history from Troy to Augustine" b
...
[8] ^ 532
    await chat("I love history and mythology' I'd like to learn something new about Greece' any suggestions?")
```

Optional Tasks

- 1. After the required notebooks are run successfully, you can explore and learn about Planner in Semantic Kernel by executing the Jupyter Notebook named 05-using-the-planner.ipynb using the Python programming language. To do this, navigate to 05-using-the-planner.ipynb and run the notebook cell-by-cell. Choose **Python <your_version>** whenever a prompt appears on top of the screen.
- 2. In the following function, modify **useAzureOpenAI** to **True**, set the values of **deployment_name**, **endpoint**, and **api_key**, which you noted earlier, and execute it.

```
import semantic_kernel as sk
from semantic_kernel.connectors.ai.open_ai import OpenAIChatCompletion, AzureChatCompletion

kernel = sk.Kernel()

useAzureOpenAI = True

# Configure AI backend used by the kernel
if useAzureOpenAI:
    deployment, api_key, endpoint = sk.azure_openai_settings_from_dot_env()
    kernel.add_chat_service("chat_completion", AzureChatCompletion(deployment_name="gpt-35-turbo", endpoint="https://aiservice.openai.azure.co
else:
    api_key, org_id = sk.openai_settings_from_dot_env()
    kernel.add_chat_service("gpt-3.5", OpenAIChatCompletion(ai_model_id="gpt-3.5-turbo", api_key=api_key, org_id=org_id))

(1) ✓ 1.7s Python
```

Run the remaining cells of code as usual.

3. After successfully completing the previous notebook, navigate to 06-memory-and-embeddings.ipynb and run the notebook cell-by-cell. Choose **Python** <your_version> whenever a prompt appears on top of the screen.
4. In the following function, modify **useAzureOpenAI** to **True**, set the values of **deployment_name**, **endpoint**, and **api_key** for both of your models, which you noted earlier, and execute it.

```
kernel = sk.Kernel()

useAzureOpenAI = True

# Configure AI service used by the kernel
if useAzureOpenAI:
    deployment, api_key, endpoint = sk.azure_openai_settings_from_dot_env()
    # next line assumes chat deployment name is "turbo", adjust the deployment name to the value of your chat model if needed
    azure_chat_service = AzureChatCompletion(deployment_name="gpt-35-turbo", endpoint="https://aiservice.openai.azure.com/", api_key=
    # next line assumes embeddings deployment name is "text-embedding", adjust the deployment name to the value of your chat model if needed
    azure_text_embedding = AzureTextEmbedding(deployment_name="text-embedding-ada-002", endpoint="https://aiservice.openai.azure.com/", api_k
    kernel.add_chat_service("chat_completion", azure_chat_service)
    kernel.add_text_embedding_generation_service("ada", azure_text_embedding)
else:
    api_key, org_id = sk.openai_settings_from_dot_env()
    oai_chat_service = OpenAIChatCompletion(ai_model_id="gpt-3.5-turbo", api_key=api_key, org_id=org_id)
    oai_text_embedding = OpenAITextEmbedding(ai_model_id="text-embedding-ada-002", api_key=api_key, org_id=org_id)
    kernel.add_chat_service("chat-gpt", oai_chat_service)
    kernel.add_text_embedding_generation_service("ada", oai_text_embedding)

kernel.register_memory_store(memory_store=sk.memory.VolatileMemoryStore())
kernel.import_skill(sk.core_skills.TextMemorySkill())

Python
```

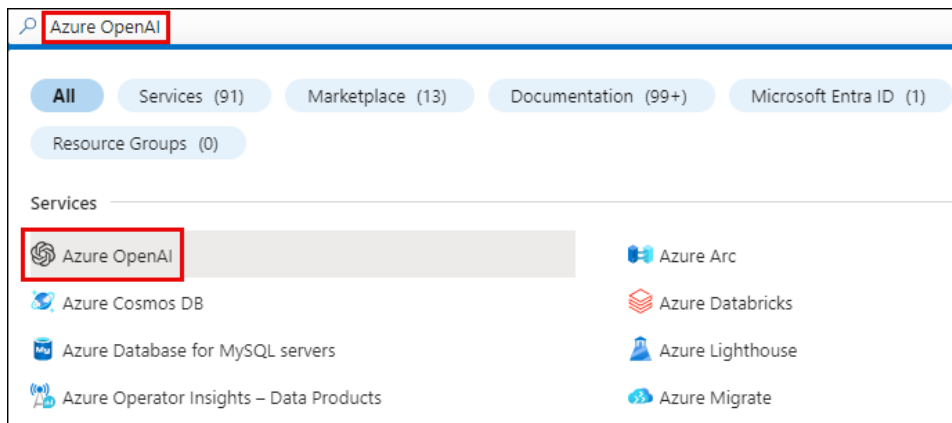
Run the remaining cells of code as usual.

Step 3: Run the Chat Copilot App Locally- Solution Guide

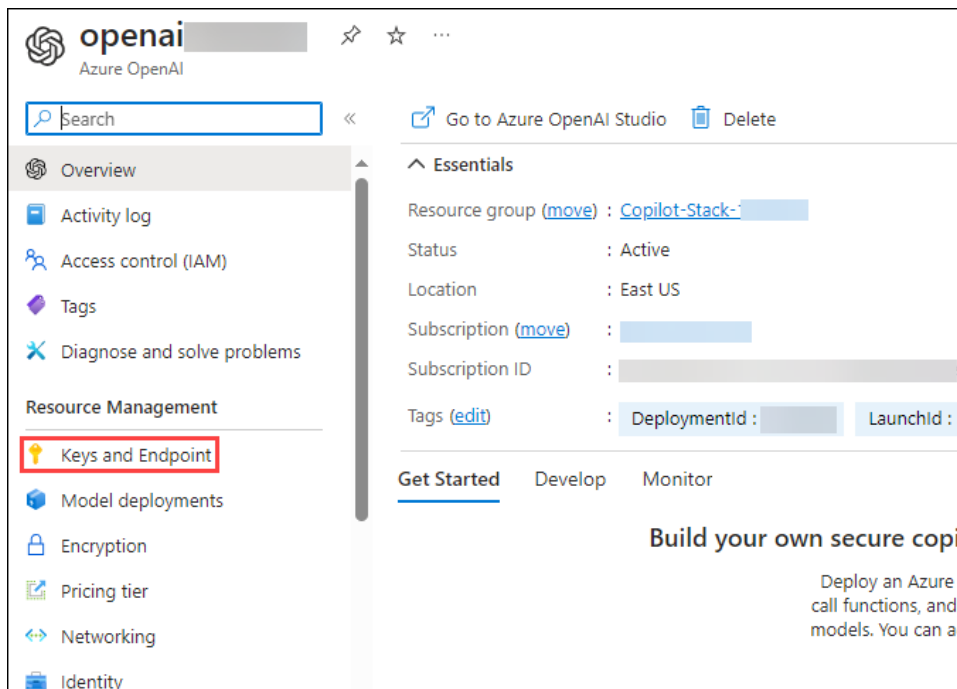
Task 1: Retrieving the Azure OpenAI Service values

From the Azure portal, you need to retrieve the Azure OpenAI Service Key, Endpoint, and LLM model names deployed in the previous Step.

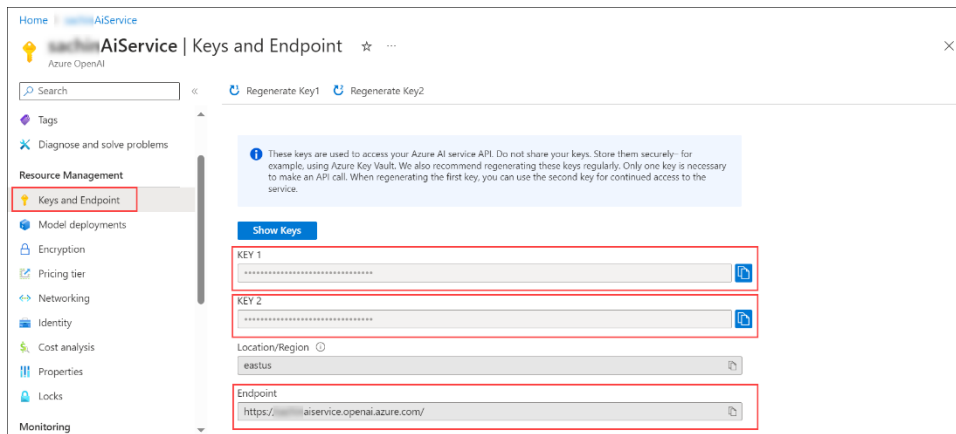
1. In the Azure Portal, search for **Azure OpenAI** and select it.



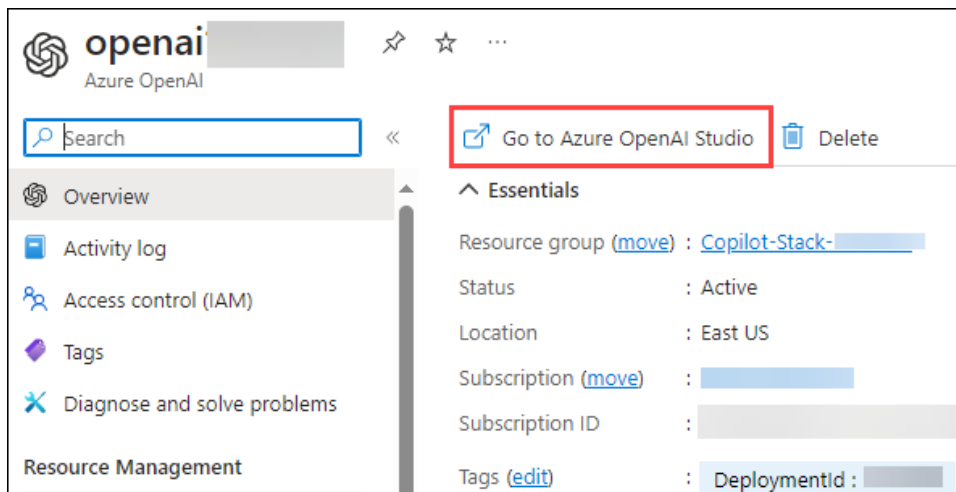
2. Select the Azure OpenAI resource created and click on **Keys and Endpoints** from the left pane.



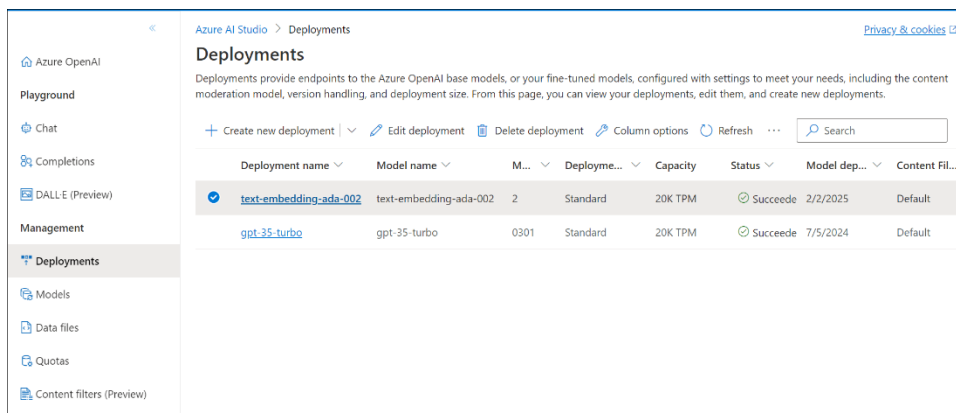
3. Copy the Keys and the Endpoint and store them in Notepad.



4. From the Overview page, click on **Go to Azure OpenAI Studio**.



5. Navigate to **Deployments** in the left navigation pane, copy the deployment names of your AI model that you deployed previously, and store them in Notepad.



Task 2: Cloning the Chat-Copilot GitHub Repo

If you have not already cloned the Chat-Copilot GitHub Repository to the environment where you're working on this lab, follow these steps to do so. Otherwise, open the cloned folder in **Visual Studio Code**.

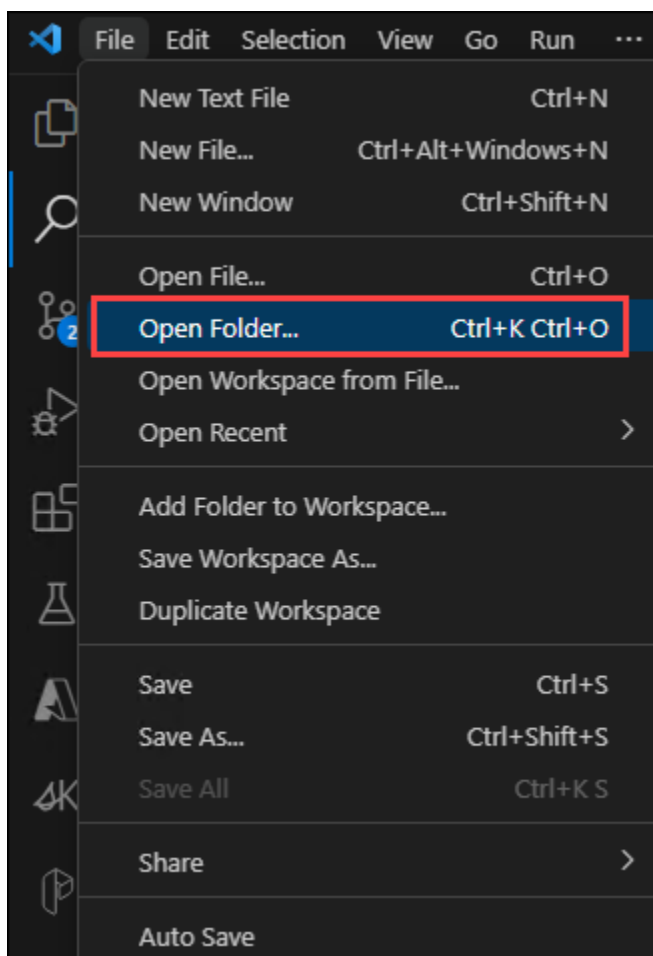
1. Open PowerShell as an administrator.
2. Navigate to the directory C:/Users/azureuser.

```
cd C:/Users/azureuser
```

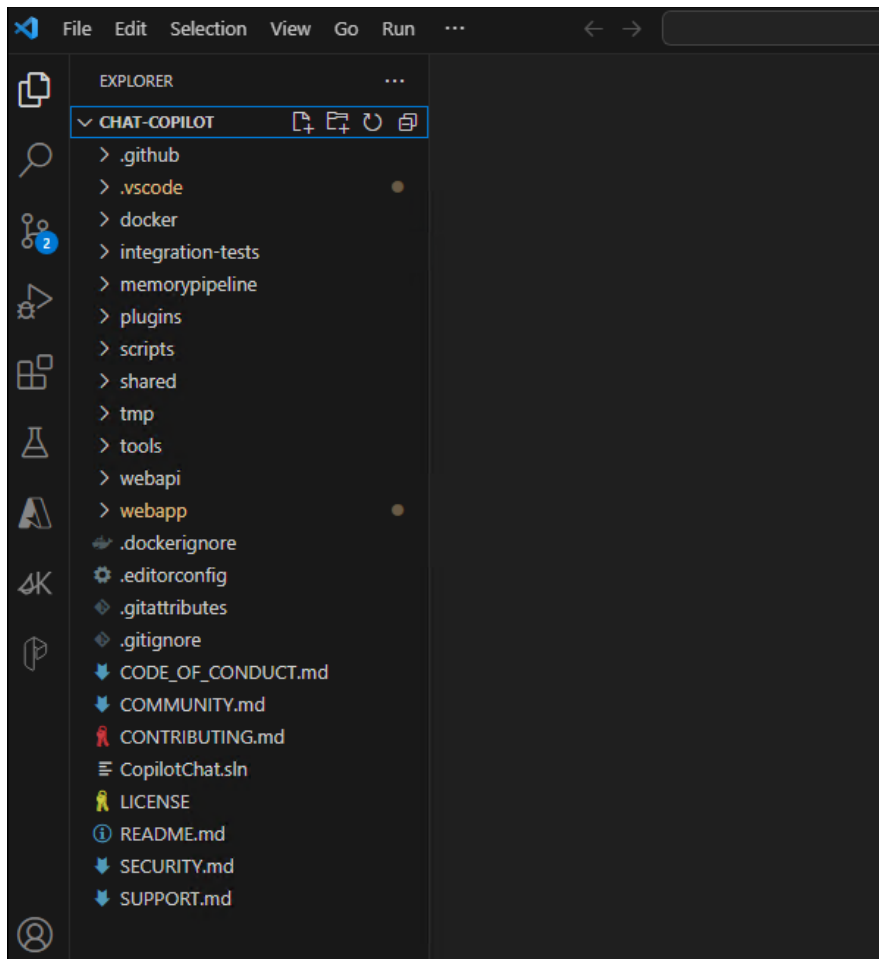
3. Clone the GitHub repository.

```
git clone https://github.com/microsoft/chat-copilot
```

4. Open Visual Studio Code and click on File> Open folder.



5. Select **CHAT-COPILOT** and review the files.



Task 3: Setting up the Environment

1. Open PowerShell as an administrator on your local machine. You need to have PowerShell Core 6+ installed, which is different from the default PowerShell installed on Windows.

2. Setup your environment by navigating to the scripts directory of chat-copilot using the command:

```
cd C:\Users\azureuser\chat-copilot\scripts\
```

3. Run the command below to install Chocolatey, dotnet-7.0-sdk, nodejs, and yarn:

```
.\Install.ps1
```

Note: If you receive an error that the script is not digitally signed or cannot execute on the system, you may need to change the execution policy or unblock the script.

Task 4: Configure and run the Chat Copilot App locally

1. Configure Chat Copilot by running the following command:

```
.\Configure.ps1 -AIService {AI_SERVICE} -APIKey {API_KEY} -Endpoint  
{AZURE_OPENAI_ENDPOINT} -CompletionModel {DEPLOYMENT_NAME} -  
EmbeddingModel {DEPLOYMENT_NAME} -PlannerModel {DEPLOYMENT_NAME}
```

Note: Provide the Azure OpenAI Service Name, Key, Endpoint, and the already deployed model names that you noted down in the previous steps.

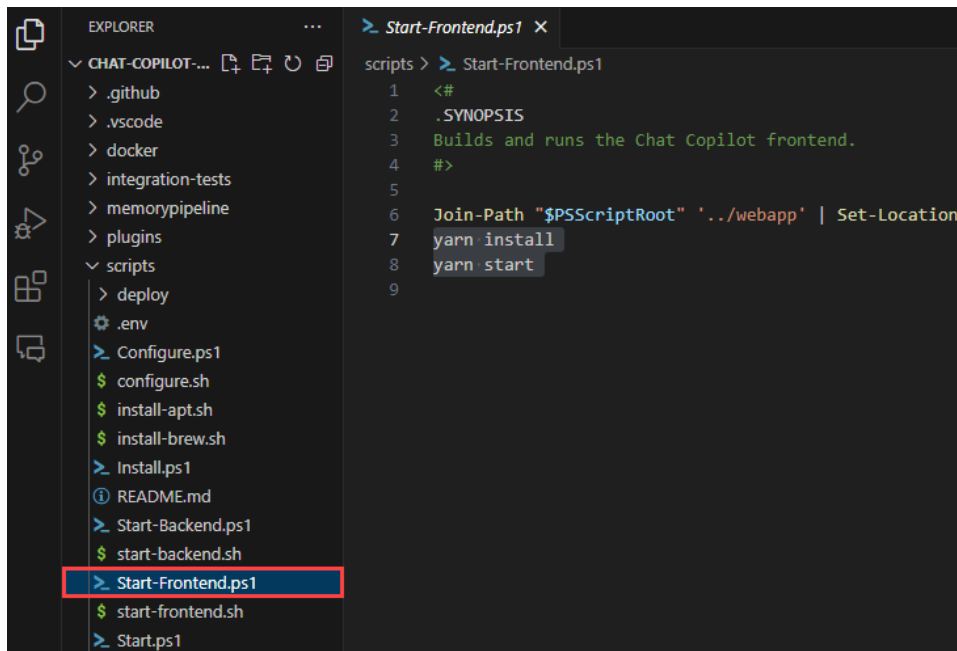
2. Finally, run Chat Copilot locally by executing the following command:

```
.\Start.ps1
```

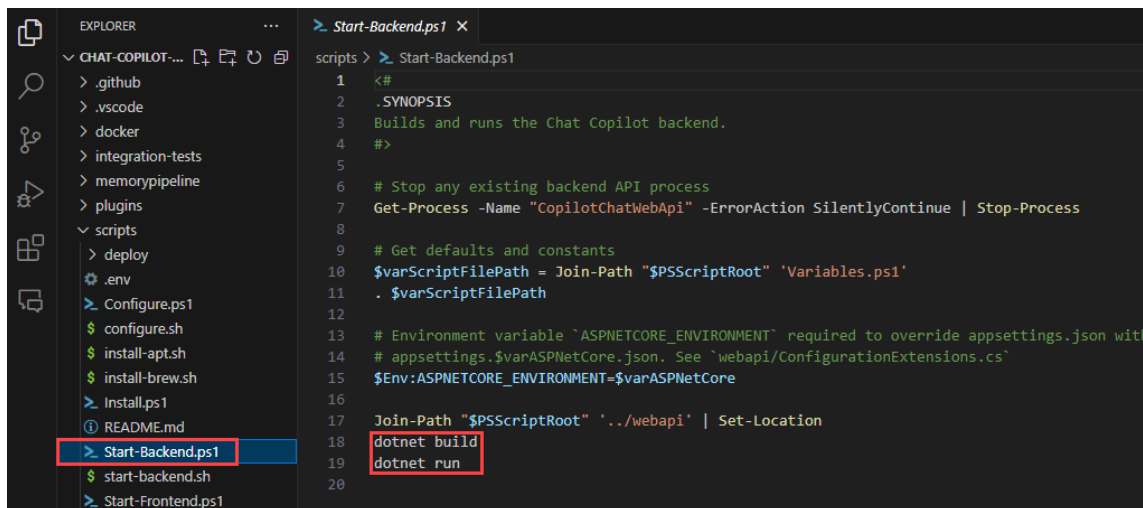
Note: This step starts both the backend API and the frontend application. It may take a few minutes for Yarn packages to install on the first run.

Note: In case of an error, follow the below steps.

3. Navigate to chat-copilot > scripts > Start-Frontend.ps1 to run the yarn commands.

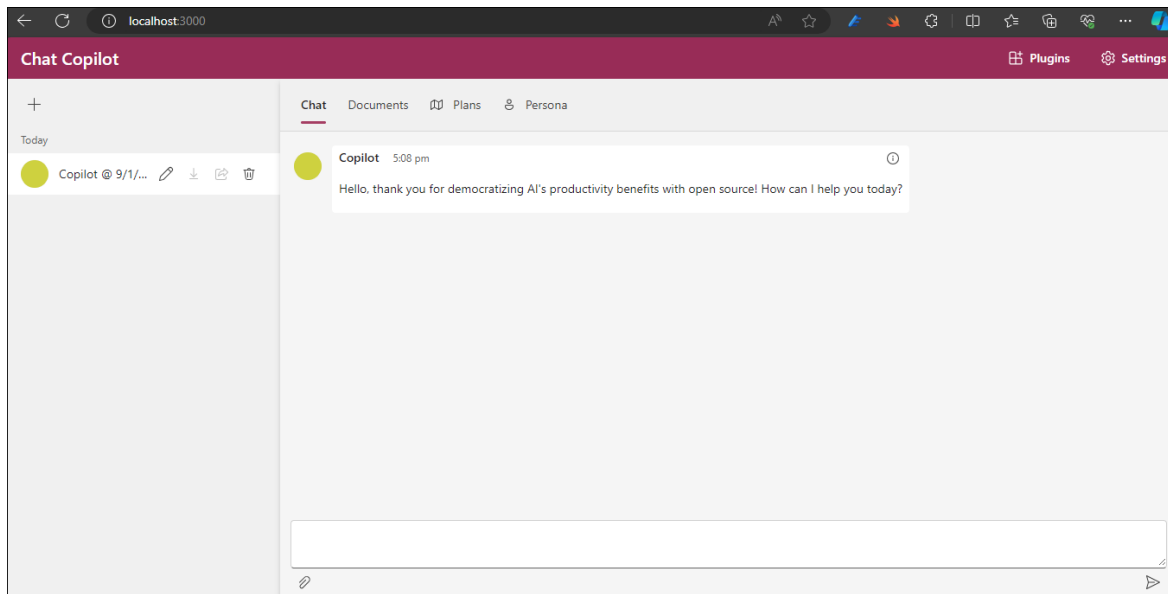


4. Navigate to chat-copilot > scripts > Start-Backend.ps1 to run the dotnet commands.



Note: Once done, navigate to the scripts directory and run the start command again.

5. You will get an output similar to this for the frontend:



6. You will get an output similar to this for the backend:

```
Select Administrator: C:\Program Files\PowerShell\pwsh.exe
MSBuild version 17.8.3+195e7f5a3 for .NET
Determining projects to restore...
All projects are up-to-date for restore.
CopilotChatShared -> C:\Users\GayatriMurali\chat-copilot-main\shared\bin\Debug\net6.0\CopilotChatShared.dll
CopilotChatWebApi -> C:\Users\GayatriMurali\chat-copilot-main\webapi\bin\Debug\net6.0\CopilotChatWebApi.dll

Build succeeded.
0 Warning(s)
0 Error(s)

Time Elapsed 00:00:06.21
info: CopilotChat.WebApi.Program[0]
      Adding plugin: Klarna Shopping.
info: CopilotChat.WebApi.Program[0]
      Added plugin: Klarna Shopping.
info: Microsoft.Hosting.Lifetime[14]
      Now listening on: https://localhost:40443
info: Microsoft.Hosting.Lifetime[0]
      Application started. Press Ctrl+C to shut down.
info: Microsoft.Hosting.Lifetime[0]
      Hosting environment: Development
info: Microsoft.Hosting.Lifetime[0]
      Content root path: C:\Users\GayatriMurali\chat-copilot-main\webapi
info: CopilotChat.WebApi.Program[0]
      Health probe: https://localhost:40443/healthz
info: Microsoft.AspNetCore.Hosting.Diagnostics[1]
      Request starting HTTP/1.1 GET https://localhost:40443/messageRelayHub - - -
info: CopilotChat.WebApi.Auth.PassThroughAuthenticationHandler[0]
      Allowing request to pass through
info: Microsoft.SemanticKernel.Connectors.AI.OpenAI.TextEmbedding.AzureOpenAITextEmbeddingGeneration[0]
      Action: GenerateEmbeddingsAsync. Azure OpenAI Deployment Name: text-embedding-ada-002.
info: Microsoft.SemanticKernel.Connectors.AI.OpenAI.TextEmbedding.AzureOpenAITextEmbeddingGeneration[0]
      Action: GenerateEmbeddingsAsync. Azure OpenAI Deployment Name: text-embedding-ada-002.
info: Microsoft.AspNetCore.Hosting.Diagnostics[1]
      Request starting HTTP/2 OPTIONS https://localhost:40443/chats/2f2b3f1e-45de-40a2-ac04-0a2db41376ea/messages - - -
info: Microsoft.AspNetCore.Hosting.Diagnostics[2]
      Request finished HTTP/2 OPTIONS https://localhost:40443/chats/2f2b3f1e-45de-40a2-ac04-0a2db41376ea/messages - 204 - - 46.8267ms
info: Microsoft.AspNetCore.Hosting.Diagnostics[1]
      Request starting HTTP/2 POST https://localhost:40443/chats/2f2b3f1e-45de-40a2-ac04-0a2db41376ea/messages - application/json 126
info: CopilotChat.WebApi.Auth.PassThroughAuthenticationHandler[0]
      Allowing request to pass through
info: Microsoft.AspNetCore.Hosting.Diagnostics[2]
      Request finished HTTP/2 POST https://localhost:40443/chats/2f2b3f1e-45de-40a2-ac04-0a2db41376ea/messages - 404 - text/plain; charset=utf-
```